

# Kubernetes Security

Chen Fei

Suneel Gandham

Ebhadulla Shaik

# What is K8S?




- K8S - **Kubernetes**
- Open source container orchestration tool
- Developed by Google by name Brog.
- Helps you manage containerized applications in different application environments

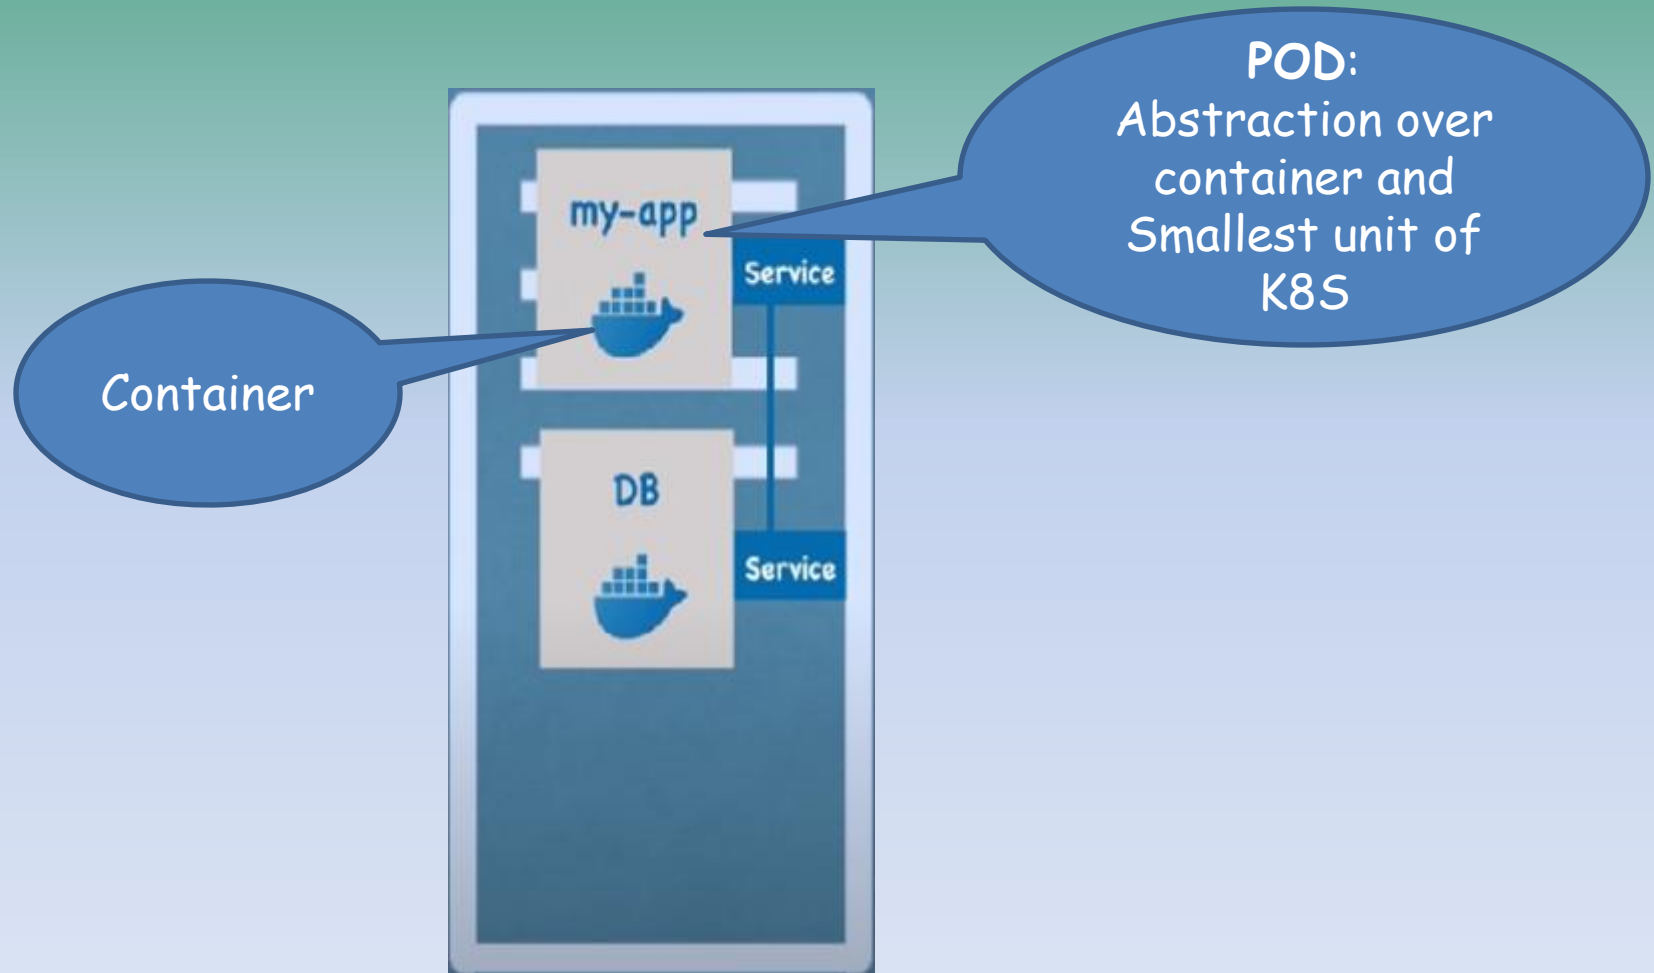
# The need for a container orchestration tool



- Trend from Monolith to Microservices
- Increased usages of containers
- Need for managing hundreds of independent microservices
- Managing them self made tools is difficult
- Availability
- Scalability
- Disaster recovery



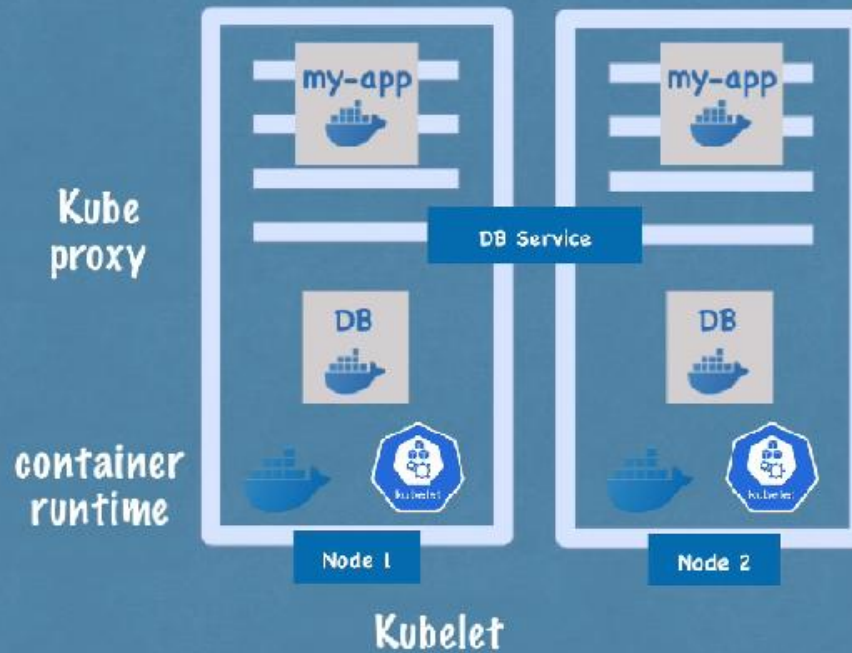
# Kubernetes Basics




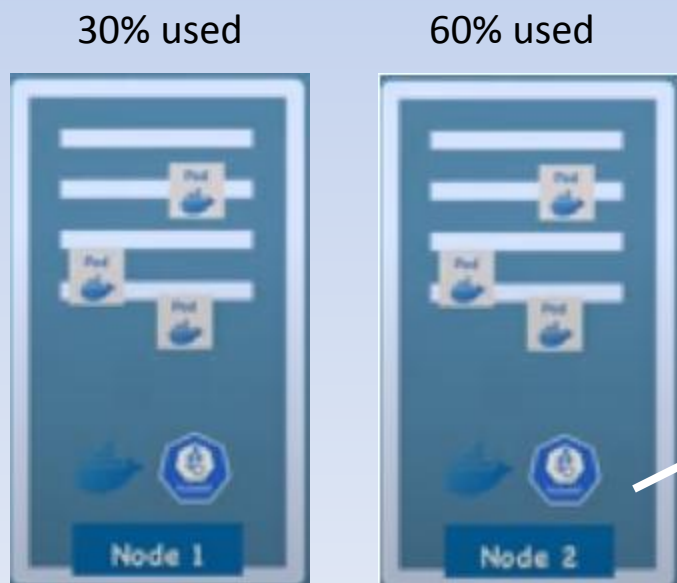
# Kubernetes Architecture




## Worker machine in K8s cluster



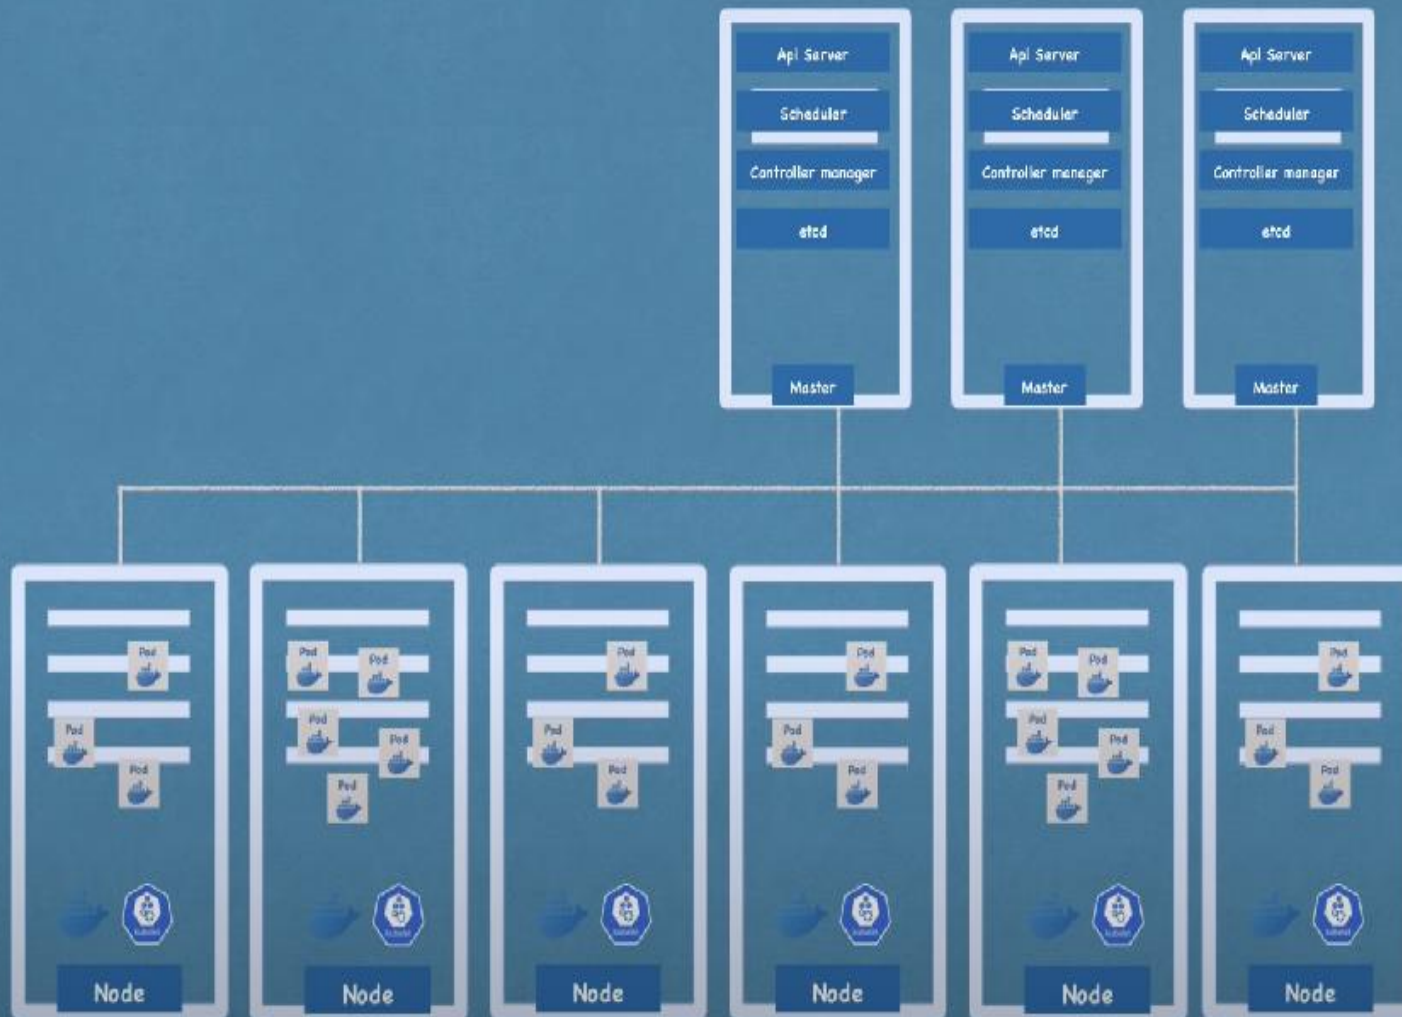
# Kubernetes Architecture



# Scale as needed



# Cluster view







# Few Key components

- Services
- Ingress
- ConfigMap
- Secrets
- Kube Proxy
- Kubelet
- Deployment
- Volumes

# How Secure Kubernetes by default?



- **Misconception:** Cloud is always secure.
- It is just that, on cloud different security tools are available which we can leverage.
- K8s main aim is container orchestration and not much focus on security.
- Don't worry... we have tools to secure the infrastructure.



# Building a secure image in CI/CD pipeline

## Causes for insecure image:

- Code from untrusted registries
- Vulnerabilities in tools of OS or code libraries

## Best Practices:

- Scanning the images in CI/CD pipeline using tools before pushing to registry.
- Scan images regularly that are pushed to registry.



# Secure Container and Pods



- Running a container with root user will make attacker's job easy to escalate the privileges.
- Attacker can easily break out from the container and can access the underlying host.





## Best Practice: Run container as non-root

- Create service user and run the container with that user

```
# create group and user
RUN groupadd -r myapp && useradd -g myapp myapp

# set ownership and permissions
RUN chown -R myapp:myapp /app

# switch to user
USER myapp
```

- Don't allow privilege escalation

```
apiVersion: v1
kind: Pod
metadata:
  name: my-app
spec:
  securityContext:
    runAsUser: 1000
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-app
spec:
  securityContext:
    allowPrivilegeEscalation: false
```



# Secure Container and Pods

- Avoid running privileged Pods
- Set limits on cluster resources for containers
  - ❖ Resource quota
  - ❖ Limit ranges for the resources

# Securing communication between Pods



- Control how the Pods communicate with each other using,

- ❖ NetworkPolicy
- ❖ Firewalls for pods.



Calico

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: my-app
spec:
  podSelector:
    matchLabels:
      name: my-app
  ingress:
    - from:
      - ipBlock:
          cidr: 172.17.0.0/16
          except:
            - 172.17.1.0/24
      - podSelector:
          matchLabels:
            name: my-db
    ports:
      - protocol: TCP
```

# Maintaining Namespaces



- Namespaces help in separating the resources, there by limiting the scope in case of attacks.






# Securing communication between Pods



- Service Mesh

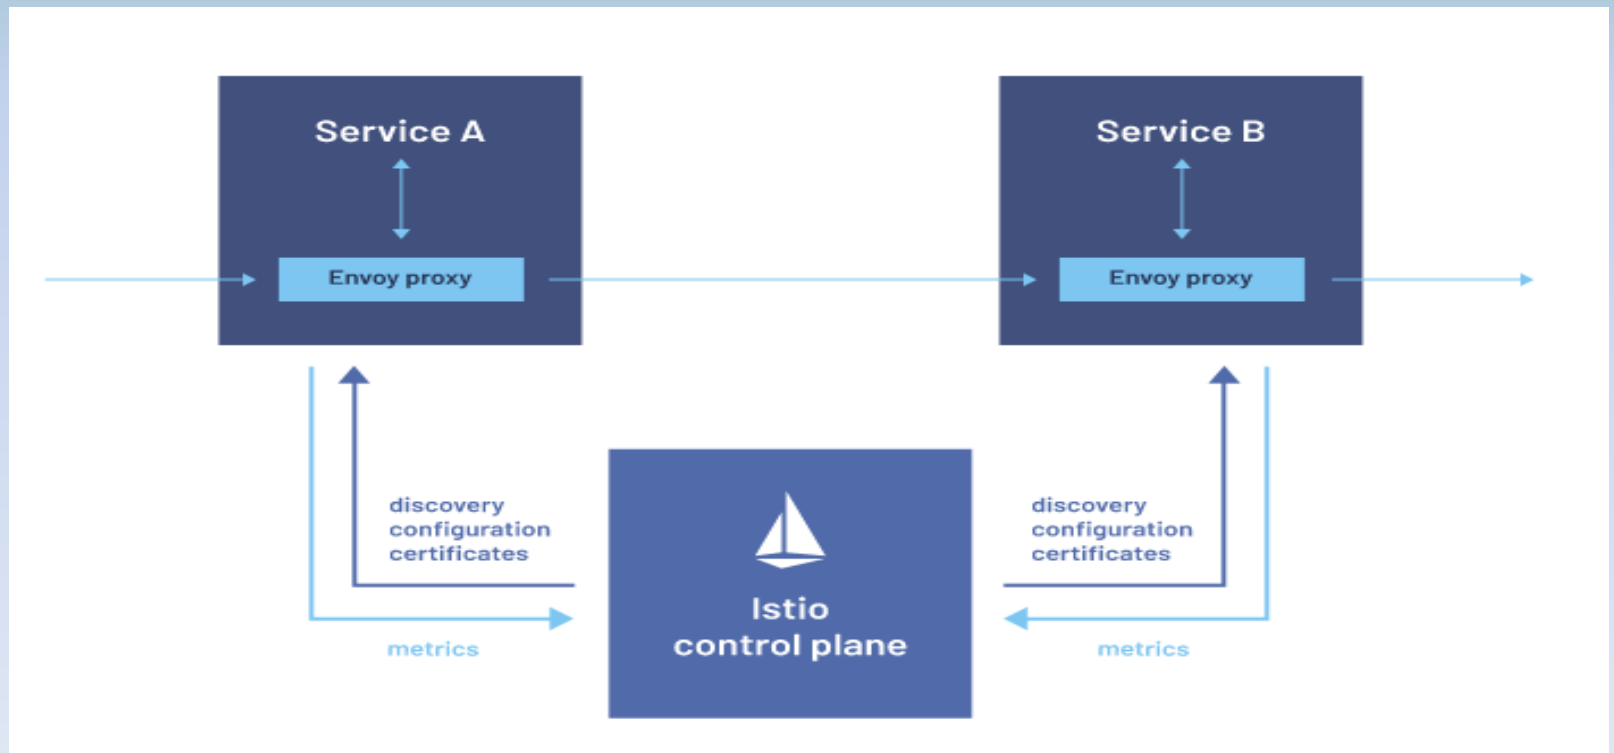
- ❖ mTLS - Mutual TLS
- ❖ Limits the communication with Network Rules
- ❖ Abstracts the rules to a layer of infrastructure and out of individual services



# Service Mesh Implementations



- Istio (native implementation) - Kiali dashboard
- Linkerd
- Consul Connect(HashiCorp)



# Securing etcd & secrets data



- By default secrets and configurations are stored as plain text
- Encoding them is not sufficient.
- Tools:
  - ❖ HashiCorp's Vault
  - ❖ Encryption tools from cloud providers
- Safe backup and recovery with K8s native K10.



# Securing access to K8s cluster



- Don't allow access application by node port.



- LoadBalancer is better



- Ingress is even better




# Securing access to K8s cluster



- Authentication

- ❖ Certificates for authenticating users
- ❖ Client certificates for human users
- ❖ Service Accounts for non-human users



# Securing access to K8s cluster



- Authorization
  - Role Based Access Control (RBAC)
  - Cluster-wide
    - ❖ Cluster Role
    - ❖ Cluster Role Binding
  - Namespace-scoped
    - ❖ Role
    - ❖ Role Binding
  - Give least possible privileges for an user

# Role Based Access Control



RoleBinding


► Attach ("Bind") a Role to a User or Group

- ✓ View/Create/Update Deployments
- ✓ View/Create/Update Services
- ✓ View/Create/Update ConfigMaps



Namespace: db

"bind"



sarah

- ✓ View/List Pods



Namespace: my-app

"bind"



tom

# Monitor and Mitigate Runtime Attacks



- Monitoring communications of Pods
- Using monitoring and logging tools
  - Kiali
  - Prometheus
  - Fluent Bit
  - Grafana



# Kiali dashboard



Overview

Graph

Applications

Workloads

Services

Istio Config

Namespace ▼ Filter by Namespace Name ▼ ↓<sup>A</sup>z

Active Filters: Namespace ✕ [Clear All Filters](#)

bookinfo

1 Label

Istio Config ✓

6 Applications ✓ 6

Healthy

- ✓ details
- ✓ kiali-traffic-generator
- ✓ mongodb
- ✓ productpage
- ✓ ratings
- ✓ reviews

istio-system

1 Label

Istio Config N/A

7 Applications ✓ 7

No traffic

Thank you